

Homework 4 — Survey of Deep Reinforcement Learning, Agentic AI, and SOTA Applications

Course: Deep Reinforcement Learning

Date: May 2026

Format: IEEE/ACM Survey Report

Abstract

This report surveys the landscape of Deep Reinforcement Learning (DRL), covering foundational algorithms, modern simulation platforms, agentic AI systems, and real-world applications across robotics, game AI, financial technology, and autonomous vehicles. We examine State-of-the-Art (SOTA) algorithms including DQN, PPO, SAC, MuZero, and Decision Transformer, and discuss emerging trends such as embodied AI, world models, and foundation model-based agents. Experimental results from five bonus experiments—DQN vs. PPO on CartPole, multi-environment PPO benchmarks, hyperparameter ablation, custom GridWorld DQN, and SAC on LunarLanderContinuous—are analyzed and contextualized within the broader literature.

Part 1 — Introduction to Deep Reinforcement Learning

1.1 Fundamental Concepts

Reinforcement Learning (RL)

Reinforcement Learning is a machine learning paradigm in which an agent learns to make sequential decisions by interacting with an environment. Unlike supervised learning, the agent receives no labeled examples; instead, it receives scalar **reward signals** that indicate the quality of its actions. The core learning signal is the mapping from states and actions to long-term cumulative reward.

The RL interaction loop: 1. The agent observes state $[s_t]$ 2. The agent selects action $[a_t]$ according to policy $[\pi]$ 3. The environment transitions to $[s_{t+1}]$ and emits reward $[r_t]$ 4. The agent updates its policy to maximize future rewards

Markov Decision Process (MDP)

Formally, RL is modeled as a **Markov Decision Process** (MDP), defined by the tuple $[(S, A, P, R, \gamma)]$:

Symbol	Meaning
$[S]$	State space
$[A]$	Action space
$[P(s' \mid s, a)]$	Transition probability
$[R(s, a)]$	Reward function
$[\gamma \in [0,1]]$	Discount factor

The **Markov property** states that the next state depends only on the current state and action, not on the history: $[P(s_{t+1} \mid s_t, a_t, \dots, s_0, a_0) = P(s_{t+1} \mid s_t, a_t)]$.

Reward Function

The reward function $[R: S \times A \rightarrow \mathbb{R}]$ maps state-action pairs to scalar signals. Reward design is arguably the most critical engineering decision in RL: - **Dense rewards**: Feedback at every step (e.g., CartPole +1 per alive step) - **Sparse rewards**: Feedback only at terminal events (e.g., MountainCar reward only when reaching the goal) - **Shaped rewards**: Engineered intermediate signals to guide learning

Policy and Value Functions

A **policy** $[\pi(a \mid s)]$ defines the agent's behavior: the probability distribution over actions given a state. The objective is to find the optimal policy $[\pi^*]$ that maximizes expected cumulative reward.

Two key value functions: - **State-value function**: $[V^\pi(s) = \mathbb{E}^\pi[\sum_{t=0}^{\infty} \gamma^t r_t \mid s_0 = s]]$ - **Action-value function (Q-function)**: $[Q^\pi(s, a) = \mathbb{E}^\pi[\sum_{t=0}^{\infty} \gamma^t r_t \mid s_0 = s, a_0 = a]]$

The **Bellman equation** provides a recursive relationship: $[Q^\pi(s, a) = R(s, a) + \gamma \mathbb{E}_{s'}[V^\pi(s')]]$

Exploration vs. Exploitation

A central challenge in RL is the **exploration-exploitation dilemma**: the agent must balance: - **Exploration**: Trying new actions to discover potentially better strategies - **Exploitation**: Leveraging known good actions to maximize immediate reward

Common strategies: $[\epsilon]$ -greedy (DQN), entropy regularization (SAC), curiosity-driven exploration (ICM), Upper Confidence Bound (UCB).

1.2 Deep Reinforcement Learning Algorithms

DQN (Deep Q-Network, Mnih et al. 2015)

DQN extends Q-learning to high-dimensional state spaces using a neural network to approximate the Q-function.

Core Idea: Approximate $[Q(s, a; \theta)]$ with a deep neural network, stabilized by: 1.

Experience Replay: Store transitions $[(s, a, r, s')]$ in a buffer; sample mini-batches randomly to break correlation 2. **Target Network:** A periodically-frozen copy of the Q-network for stable Bellman targets

Loss:
$$\mathcal{L}(\theta) = \mathbb{E} \left[(r + \gamma \max_{a'} Q(s', a'; \theta^-) - Q(s, a; \theta))^2 \right]$$

Dimension	Description
Advantages	Simple, well-understood, works on Atari
Weaknesses	Discrete actions only; overestimates Q-values; sample inefficient
Applications	Atari games, discrete control tasks

Double DQN (van Hasselt et al. 2016)

Standard DQN suffers from **Q-value overestimation** because the same network selects and evaluates actions. Double DQN decouples these:

$$y = r + \gamma Q(s', \arg \max_{a'} Q(s', a; \theta); \theta^-)$$

This separates action selection (online network $[\theta]$) from value estimation (target network $[\theta^-]$), significantly reducing overestimation bias and improving stability.

PPO (Proximal Policy Optimization, Schulman et al. 2017)

PPO is an **on-policy** policy gradient method that constrains each update to prevent destructive large policy changes.

Clipped Surrogate Objective:
$$\mathcal{L}^{\text{CLIP}}(\theta) = \mathbb{E}_t \left[\min \left(r_t(\theta) \hat{A}_t, \text{clip}(r_t(\theta), 1 - \epsilon, 1 + \epsilon) \hat{A}_t \right) \right]$$

where $[r_t(\theta) = \frac{\pi_{\theta}(a_t \mid s_t)}{\pi_{\theta_{\text{old}}}(a_t \mid s_t)}]$ is the probability ratio and $[\hat{A}_t]$ is the advantage estimate.

Dimension	Description
Advantages	Stable; simple implementation; RLHF industry standard
Weaknesses	On-policy: poor sample efficiency; does not reuse past data
Applications	ChatGPT/Claude alignment (RLHF), robotics, game AI

Our experiments confirm PPO achieves perfect CartPole-v1 score (500) within ~150 episodes, substantially faster than DQN (which catastrophically forgets after ~300 episodes).

A2C / A3C (Mnih et al. 2016)

Advantage Actor-Critic (A2C) uses two networks: an **actor** $[\pi_{\theta}(a \mid s)]$ and a **critic** $[V_{\phi}(s)]$. The advantage $[A(s, a) = Q(s, a) - V(s)]$ reduces variance in policy gradient estimates.

Asynchronous Advantage Actor-Critic (A3C) runs multiple agents in parallel environments, each asynchronously updating a global network. A2C is the synchronous variant that waits for all workers before updating.

Dimension	Description
Advantages	More stable than vanilla policy gradient; parallelizable
Weaknesses	Higher variance than PPO; synchronization overhead in A3C
Applications	Continuous control, game playing

SAC (Soft Actor-Critic, Haarnoja et al. 2018)

SAC is an **off-policy** actor-critic method for continuous action spaces, built on the **maximum entropy** framework.

Objective: $[\pi^* = \arg\max_{\pi} \mathbb{E}[\sum_t r(s_t, a_t) + \alpha \mathcal{H}(\pi(\cdot \mid s_t))]]$

where $[\mathcal{H}(\pi) = -\mathbb{E}[\log \pi(a \mid s)]]$ is the entropy term and $[\alpha]$ is a temperature parameter that is **automatically tuned** during training.

Dimension	Description
Advantages	SOTA on continuous control; off-policy (sample efficient); automatic entropy tuning
Weaknesses	More complex implementation; slower per-step than PPO
Applications	Robot manipulation, locomotion, LunarLander continuous

Our Bonus 5 experiment achieves 271.7 ± 15.5 on LunarLanderContinuous-v3 (well above the 200 success threshold), using 300k training steps with `net_arch=[256,256]`. The entropy coefficient α visibly auto-tunes from ~ 1.0 to ~ 0.2 over training, confirming SAC's maximum-entropy framework in action. This demonstrates SAC's superiority for continuous control tasks that DQN fundamentally cannot handle.

TD3 (Twin Delayed DDPG, Fujimoto et al. 2018)

TD3 addresses overestimation in DDPG through three techniques: 1. **Twin Critics**: Use two Q-networks, take the minimum for target computation 2. **Delayed Policy Updates**: Update actor less frequently than critics 3. **Target Policy Smoothing**: Add noise to target actions

Dimension	Description
Advantages	More stable than DDPG; competitive with SAC on many benchmarks
Weaknesses	Requires more careful hyperparameter tuning than SAC
Applications	Continuous robot control, MuJoCo benchmarks

MuZero (Schrittwieser et al. 2020)

MuZero learns a **latent dynamics model** without requiring knowledge of environment rules. It combines model-based planning (MCTS) with model-free value estimation, achieving superhuman performance on Atari, Chess, Go, and Shogi.

Three learned components: - **Representation function**: $[h_{\theta}(o_t) \rightarrow s_t]$ (encode observations) - **Dynamics function**: $[g_{\theta}(s_t, a_t) \rightarrow (r_t, s_{t+1})]$ (predict transitions) - **Prediction function**: $[f_{\theta}(s_t) \rightarrow (\pi_t, v_t)]$ (policy and value)

Dimension	Description
Advantages	Works without environment model; strong planning ability
Weaknesses	High computational cost; complex training
Applications	Board games, Atari, video compression optimization

Decision Transformer (Chen et al. 2021)

Decision Transformer reformulates RL as a **sequence modeling** problem. It uses a GPT-style causal transformer to predict actions conditioned on desired return, past states, and past actions:

$[(R_1, s_1, a_1, R_2, s_2, a_2, \dots, R_t, s_t, a_t)]$

The model learns to output the action $[a_t]$ that achieves the specified return-to-go $[R_t]$.

Dimension	Description
Advantages	Leverages transformer scalability; works well offline
Weaknesses	Cannot improve beyond dataset quality; no online adaptation
Applications	Offline RL, robot learning from demonstrations

Offline RL

Offline RL (batch RL) learns policies purely from pre-collected datasets, without any environment interaction. Key challenge: **distributional shift** — the learned policy may visit states not covered by the dataset, leading to catastrophic Q-value overestimation.

Key methods: - **BCQ** (Batch-Constrained Q-learning): Constrain policy to stay near dataset actions - **CQL** (Conservative Q-Learning): Penalize Q-values for out-of-distribution actions - **IQL** (Implicit Q-Learning): Avoid querying OOD actions entirely

Applications: Healthcare (cannot do online trials), autonomous driving (safety constraints).

Hierarchical RL

Hierarchical RL decomposes long-horizon tasks into sub-goals: - **High-level policy**: Sets sub-goals (e.g., "go to room A") - **Low-level policy**: Executes primitive actions to reach sub-goals

Key frameworks: **HIRO** (Hierarchical Reinforcement Learning with Off-policy Correction), **Options Framework**, **Feudal Networks**.

Applications: Multi-room navigation, robot arm pick-and-place with multiple objects, long-horizon planning.

Part 2 — Survey of DRL Systems and Platforms

2.1 NVIDIA Isaac Gym / Isaac Sim

Architecture: GPU-accelerated physics simulation running thousands of parallel environments simultaneously on a single GPU. Isaac Gym uses NVIDIA PhysX and provides direct GPU tensor APIs eliminating CPU-GPU data transfer bottlenecks.

Key Features: - Runs 4,096–8,192 parallel environments per GPU - Direct PyTorch tensor integration - Supports rigid body, articulated body, and soft body simulation

Algorithms Supported: PPO, SAC, TD3, AMP (Adversarial Motion Prior)

Applications: Humanoid locomotion (Unitree G1, H1), dexterous manipulation, sim-to-real transfer

Strengths: Unprecedented simulation throughput; essential for sample-hungry RL algorithms

Limitations: NVIDIA GPU required; steep learning curve; limited rendering fidelity vs. Isaac Sim

2.2 Habitat Lab (Meta AI)

Architecture: A modular embodied AI research platform for training agents in photorealistic 3D environments (Matterport3D, Gibson, HM3D datasets).

Supported Tasks: PointGoal Navigation, ObjectGoal Navigation, Room Rearrangement, Social Navigation

Algorithms: PPO, DD-PPO (Decentralized Distributed PPO with 64–512 GPUs)

Strengths: Large-scale indoor scene datasets; strong sim-to-real transfer results

Limitations: Primarily navigation-focused; requires significant compute for large-scale training

2.3 CARLA (Open-source Autonomous Driving Simulator)

Architecture: Unreal Engine 4-based autonomous driving simulator with sensor suite (RGB cameras, LiDAR, radar, GPS), traffic management, and weather simulation.

Supported Algorithms: Any RL framework; commonly used with PPO, SAC, DDPG for end-to-end driving

Applications: Waypoint following, obstacle avoidance, multi-agent traffic simulation

Strengths: High visual fidelity; comprehensive sensor suite; active community

Limitations: High computational cost per environment step; limited physics accuracy vs. specialized simulators

2.4 Unity ML-Agents

Architecture: Unity game engine extended with Python API for RL training. Environments can be built with Unity's visual editor and scripted behaviors.

Algorithms: PPO, SAC (built-in), custom algorithms via Python API

Strengths: Artist-friendly environment creation; cross-platform; large ecosystem

Limitations: Lower physics accuracy than Isaac Gym; slower than specialized simulators

2.5 RLlib (Ray Framework)

Architecture: Distributed RL library built on Apache Ray. Supports multi-GPU and multi-node distributed training with minimal code changes.

Algorithms: PPO, SAC, DQN, TD3, APPO, IMPALA, MADDPG, and 30+ others

Strengths: Production-scale distributed training; multi-agent support; easy scaling

Limitations: Complex configuration; heavier than SB3 for single-machine research

2.6 Stable-Baselines3 (SB3)

Architecture: Clean, well-documented implementations of RL algorithms in PyTorch. Designed for research reproducibility and benchmarking.

Algorithms: DQN, PPO, SAC, TD3, A2C, DDPG, HER, TQC (via sb3-contrib)

Strengths: Beginner-friendly; reliable implementations; active maintenance; strong community

Limitations: Single-machine only; no distributed training; limited custom environment integration

Used in all five bonus experiments in this homework.

2.7 FinRL

Architecture: Multi-agent RL framework specifically designed for financial applications, built on top of Stable-Baselines3 and RLlib.

Features: Integrates financial data APIs (Yahoo Finance, Alpaca), portfolio optimization environments, risk metrics

Applications: Stock trading, portfolio management, cryptocurrency trading

2.8 MineDojo

Architecture: Minecraft-based open-world learning platform with 1,000+ diverse tasks defined via natural language and video demonstrations (YouTube gameplay).

Key Contribution: Foundation models for embodied agents (GROOT, VPT — Video PreTraining)

Strengths: Open-ended learning; rich natural language task specification; huge internet-scale training data

2.9 OpenVLA

Architecture: Vision-Language-Action model for robot manipulation, trained on Open X-Embodiment dataset (70+ robot types, 1M+ demonstrations). Fine-tunable for specific robot setups.

Applications: Robot manipulation from visual observation and natural language instruction

2.10 AirSim (Microsoft)

Architecture: Unreal Engine-based photorealistic simulator for drones and ground vehicles with physics-accurate dynamics and sensor simulation.

Applications: UAV navigation, computer vision research, autonomous vehicle testing

Part 3 — Agentic AI and Autonomous Agents

3.1 What is Agentic AI?

Agentic AI refers to AI systems that autonomously pursue goals over extended time horizons, taking sequences of actions that include planning, tool use, memory retrieval, and self-correction. Unlike single-inference LLMs, agentic systems exhibit:

- **Goal persistence:** Maintain objectives across many steps
- **Tool use:** Invoke external APIs, databases, code interpreters, browsers
- **Memory:** Episodic (conversation history), semantic (knowledge retrieval), procedural (learned skills)
- **Planning:** Decompose complex goals into sub-tasks
- **Reflection:** Self-critique and iteratively improve outputs

3.2 RLHF (Reinforcement Learning from Human Feedback)

RLHF is the technique that enabled ChatGPT, Claude, and Gemini to follow human instructions. The pipeline:

1. **Supervised Fine-Tuning (SFT):** Train a base LLM on high-quality demonstrations
2. **Reward Model Training:** Train a Bradley-Terry reward model from human pairwise preference comparisons
3. **PPO Fine-Tuning:** Use PPO to optimize the LLM against the reward model, with a KL penalty to prevent reward hacking

RLHF Key Challenge: Reward hacking — the model learns to game the reward model rather than genuinely improving.

3.3 RLAIF (Reinforcement Learning from AI Feedback)

RLAIF (Bai et al. 2022, "Constitutional AI") replaces human annotators with an AI judge:

1. A **Constitutional AI** model critiques and revises its own responses according to a set of principles
2. AI-generated preference labels are used to train the reward model (replacing expensive human annotation)
3. PPO fine-tunes the final model

Advantages over RLHF: Scales without human annotation cost; more consistent; enables continuous self-improvement

3.4 Tool-Using Agents

Modern LLM agents augmented with tools: **ReAct** (Reason + Act) framework interleaves reasoning traces with tool calls. Tools include: - Web search (Bing, Google) - Code execution (Python interpreter) - File system access - Database queries - External APIs

OpenAI Agents API (2025): Production-grade agent infrastructure with handoffs, guardrails, tool use, and streaming.

3.5 Autonomous Planning

Chain-of-Thought (CoT) prompting enables LLMs to decompose tasks into intermediate steps. **Tree of Thoughts (ToT)** extends this to tree-structured search over reasoning paths.

For robotics, **SayCan** (Google, 2022) grounds LLM plans in robot affordances: the LLM proposes high-level steps while a value function scores their physical feasibility.

3.6 Multi-Agent Systems

Multi-agent RL (MARL) trains multiple interacting agents: - **Cooperative:** Agents share reward (StarCraft multi-agent challenge, SMAC) - **Competitive:** Zero-sum (AlphaStar, OpenAI Five) - **Mixed:** Social dilemmas (traffic optimization, market simulation)

Key Challenge: Non-stationarity — each agent's environment changes as other agents learn, violating the Markov assumption.

3.7 Memory Systems in Agents

Memory Type	Mechanism	Example
Working memory	In-context window	LLM conversation history
Episodic memory	Vector database retrieval	RAG systems
Semantic memory	Knowledge graph / embeddings	Entity knowledge

Memory Type	Mechanism	Example
Procedural memory	Fine-tuning / LoRA	Skill acquisition

3.8 Representative Agentic Systems

Voyager (Wang et al. 2023): Lifelong learning agent in Minecraft using GPT-4 for skill library construction, automatic curriculum generation, and iterative code-based skill writing. Demonstrates progressive skill acquisition without manual reward engineering.

AutoGPT: Early autonomous agent that loops LLM calls with tool use (web search, file I/O, code execution). Demonstrated long-horizon task completion but prone to error compounding.

OpenAI Deep Research: Agentic system for multi-step web research, generating comprehensive reports by autonomously querying, reading, and synthesizing web sources.

RT-2 (Robotic Transformer 2, Google 2023): Vision-language-action model that directly generates robot actions (joint torques) from visual observations and natural language instructions, leveraging internet-scale pretraining.

3.9 DRL's Role in Agentic AI

DRL enables agentic systems to:

1. **Learn reward models** from preference data (RLHF/RLAIF)
2. **Plan in latent spaces** using world models (MuZero, DreamerV3)
3. **Control physical robots** with continuous action policies (SAC, TD3)
4. **Optimize multi-step decision-making** beyond what greedy strategies achieve

Future AGI Directions: World model-based planning (like DreamerV3), combined with LLM reasoning and tool use, represents a plausible path toward more general autonomous agents.

Part 4 — DRL Applications in Different Research Domains

4.1 Robotics

Problem Domain

Robotics demands real-time, high-precision continuous control in unstructured environments. Key challenges:

- High-dimensional continuous action spaces (joint torques, velocities)
- Partial observability from sensors
- Sim-to-Real gap (trained in simulation, deployed on physical hardware)
- Safety constraints during learning

Why DRL?

Traditional control (PID, model predictive control) requires precise system models. DRL learns controllers directly from experience, adapting to unknown dynamics and enabling skills that are difficult to model analytically (e.g., dexterous manipulation, bipedal locomotion over rough terrain).

SOTA Approaches

Locomotion: NVIDIA Isaac Lab trains humanoid robots (Unitree G1) using massive parallelism (8,192 environments) with PPO + domain randomization. ETH Zurich's ANYmal quadruped achieves robust locomotion on stairs and challenging terrain using PPO with teacher-student distillation.

Manipulation: OpenVLA and RT-2 demonstrate that pretrained vision-language models can be fine-tuned for robot manipulation, enabling generalization to novel objects and instructions. Diffusion Policy (Chi et al. 2023) uses diffusion models to represent multimodal action distributions, achieving state-of-the-art on push-T and block stacking tasks.

Sim-to-Real Transfer: Domain Randomization (varying physics parameters during training) is the dominant technique. Recent work adds **Adaptive Domain Randomization (ADR)** that automatically adjusts randomization range based on training progress.

Key Datasets and Simulators

- **Isaac Gym/Lab:** GPU-parallelized physics simulation
- **Open X-Embodiment:** 1M+ demonstrations from 70+ robot types
- **RoboMimic:** Manipulation demonstrations with multiple operators
- **MuJoCo:** Standard locomotion benchmarks (HalfCheetah, Ant, Hopper)

Future Directions

Foundation models for robotics (similar to GPT for language), combining internet-scale pretraining with robot-specific fine-tuning. Dexterous hand manipulation remains an open challenge.

4.2 Game AI

Problem Domain

Games provide controlled environments with well-defined rules, clear objectives, and reproducible evaluation. They range from simple grid worlds to complex real-time strategy games with partial observability and long-horizon planning.

Why DRL?

Games are ideal testbeds because: (1) fast simulation, (2) clear reward signals, (3) no safety constraints, (4) easily parallelizable. Breakthroughs in games have consistently preceded real-world applications.

SOTA Approaches

AlphaGo → AlphaZero → MuZero: DeepMind's progression from Go-specific (AlphaGo, 2016) to domain-agnostic (AlphaZero, 2018) to model-learning (MuZero, 2020) represents the arc from narrow to general game-playing AI.

AlphaStar (Vinyals et al. 2019): Defeated professional StarCraft II players using a combination of supervised learning from human replays, self-play via league training, and PPO with LSTM for handling partial observability and long time horizons (up to 16,000 action steps per game).

OpenAI Five (2019): Five coordinated PPO agents defeated world champion Dota 2 players. Key insight: massive scale (180 years of self-play per day) enables emergent team coordination without hand-coded cooperation.

Minecraft (VPT, OpenAI 2022): Video PreTraining learned from 70,000 hours of YouTube gameplay using inverse dynamics models, enabling complex behaviors (crafting, building, combat) in open-ended environments.

Voyager (2023): GPT-4 as an autonomous Minecraft agent that writes, stores, and reuses code-based skills, achieving progressively harder milestones without human reward engineering.

Key Insights for the Field

1. **Scale matters:** More compute + more self-play consistently improves performance
2. **Curriculum learning:** Progressive task difficulty crucial for complex games
3. **Self-play:** Enables open-ended skill development without human-labeled data

Future Directions

Real-time strategy games with even larger state spaces (Full StarCraft map without fog-of-war), generalist game-playing agents across multiple game types, and using game AI insights for real-world decision-making.

4.3 FinTech (Financial Technology)

Problem Domain

Financial markets are sequential decision problems with: - High-dimensional, noisy, non-stationary state spaces (price series, fundamental data, market microstructure) - Delayed and sparse rewards (profit realized over days/weeks) - Partial observability (hidden order flow, insider information) - High stakes and regulatory constraints

Why DRL?

Traditional quantitative finance uses handcrafted features and statistical models (mean-reversion, momentum). DRL can discover complex non-linear relationships and adapt to changing market regimes automatically.

SOTA Approaches

Algorithmic Trading: FinRL (Liu et al. 2022) benchmarks DQN, PPO, A2C, TD3, and SAC on stock trading tasks using Yahoo Finance data. SAC consistently outperforms others due to its entropy-regularized exploration, which prevents over-fitting to specific market regimes.

Portfolio Optimization: Deep Portfolio Management (Jiang et al. 2017) uses a CNN to process price series and output portfolio weights. EIIE (Ensemble of Identical Independent Evaluators) architecture enables direct rebalancing without transaction cost approximation.

Market Making: Two-sided market-making agents (bid/ask spread management) use DQN with inventory penalties. Avellaneda-Stoikov-inspired reward functions balance profit with inventory risk.

Risk Management: DRL-based VaR (Value at Risk) optimization treats risk constraints as part of the MDP reward structure, enabling dynamic hedging strategies that adapt to volatility regimes.

Challenges and Limitations

1. **Non-stationarity:** Market regimes shift (bull/bear markets, volatility clusters)
2. **Survivorship bias:** Historical datasets exclude delisted stocks
3. **Transaction costs:** Naive DRL ignores slippage and market impact
4. **Overfitting:** Financial DRL models often overfit to in-sample data

Key Datasets

- Yahoo Finance, Alpaca Markets API (free)
- CRSP, Compustat (academic, paid)
- Crypto: Binance, Coinbase APIs (free, high-frequency)

Future Directions

Multi-agent market simulation to study systemic risk, foundation models for financial time series, and RL-based optimal order execution for institutional trading.

4.4 Autonomous Vehicles

Problem Domain

Self-driving vehicles must navigate complex, dynamic environments while ensuring safety. Key challenges: - Real-time decision-making under uncertainty - Multi-agent interaction with human drivers - Rare but critical edge cases (emergency vehicles, unusual weather) - Strict safety requirements (failure = injury or death)

Why DRL?

Rule-based systems struggle with edge cases and require extensive manual engineering. DRL can learn adaptive policies from experience, handle complex multi-agent interactions, and generalize across diverse scenarios.

SOTA Approaches

End-to-End Driving: CARLA Challenge 2024 leaders use transformer-based architectures taking camera/LiDAR inputs and outputting steering, throttle, and brake. **TransFuser** (Chitta et al. 2022) fuses image and LiDAR features via transformer cross-attention, achieving state-of-the-art on CARLA benchmarks.

Path Planning with RL: Highway-env (Leurent, 2018) is a lightweight highway driving simulator. PPO and DQN agents learn lane-changing and merging behaviors. SAC achieves smooth, human-like driving in continuous action settings.

WAYMO's ML-Planner: Uses a learned motion forecasting model combined with a rule-based planner, with RL for long-tail scenario handling. Not fully end-to-end but uses DRL for subcomponents.

Multi-Agent Driving: SMARTS (Huawei) and **MetaDrive** (Li et al. 2022) simulate dense urban traffic for multi-agent interaction research. MARL enables agents to learn cooperative and competitive driving behaviors.

Safety Considerations

Constrained MDP (CMDP) frameworks add safety constraints: $\mathbb{E} \left[\sum_t c_t \right] \leq d$ where c_t are cost signals (e.g., collision penalty). **Safe RL** methods (CPO,

PPO-Lagrangian) optimize reward while satisfying safety constraints during training.

Key Simulators

Simulator	Focus	Fidelity
CARLA	Urban driving	High
SUMO	Traffic flow	Medium
Highway-env	Highway scenarios	Low (fast)
MetaDrive	Procedural generation	Medium
nuPlan	Real data replay	High

Future Directions

World model-based planning for autonomous driving (MILE, DriveDreamer), foundation models trained on internet-scale driving video, and sim-to-real transfer using domain randomization for rare weather and edge cases.

Part 5 — SOTA Research Trends (2025–2026)

5.1 Embodied AI

What it is: AI systems that learn through physical interaction with environments, grounding language and reasoning in sensorimotor experience. The "body" provides the interface between computation and world.

Why it's important: Pure language models lack grounding — they process symbols without understanding physical consequences. Embodied AI closes the loop between perception, action, and learning.

Current SOTA: OpenVLA, RT-2, [pi_0] (Physical Intelligence, 2024) — all demonstrate foundation models that generalize manipulation skills across robot morphologies.

Limitations: Sim-to-real gap remains significant; data collection is expensive; real-world deployment safety is unresolved.

5.2 World Models

What it is: Learned environment models that enable agents to plan in imagination (latent-space rollouts) without real-world interaction.

Key Systems: - **DreamerV3** (Hafner et al. 2023): A single model that achieves human-level performance on 150+ tasks across Atari, DMLab, MuJoCo, Minecraft using only world-model-based imagination. - **GAIA-1** (Wayve, 2023): Automotive world model that generates diverse driving scenarios from text and action conditioning for training self-driving agents. - **UniSim** (Yang et al. 2023): Universal simulator trained on diverse robot and game data, enabling zero-shot policy synthesis.

Why it's important: Reduces real-world data needs by orders of magnitude; enables planning that traditional model-free RL cannot.

5.3 Diffusion Policy

What it is: Using diffusion models (score-based generative models) to represent robot action distributions, enabling multimodal policy learning.

Key Paper: Chi et al. 2023 — Diffusion Policy outperforms behavior cloning, GAIL, and IQL on 11 robot tasks by modeling the full action distribution rather than a single deterministic output.

Why it's important: Robot tasks often have multiple valid action modes; unimodal policy representations (standard actor networks) average over modes and fail. Diffusion policy captures all valid solutions.

5.4 Multi-Agent RL (MARL) — 2025 Trends

MARL has shifted from small-scale homogeneous agents to large-scale heterogeneous systems: - **LLM-based multi-agent frameworks** (AutoGen, CrewAI, Swarm): LLM agents as actors in a MARL-like loop - **Emergent communication:** Agents develop communication protocols without supervision - **Mean-field RL:** Scalable MARL for thousands of agents (traffic, crowd simulation)

5.5 Offline RL + Foundation Models

Trend: Pre-train large offline RL models on diverse datasets, then fine-tune online with minimal real-world interaction. Analogous to GPT pretraining + RLHF fine-tuning.

Key Work: **Gato** (Reed et al. 2022): Single transformer trained to play Atari, caption images, answer questions, and control robots — same weights, task specified by context.

5.6 RL + Transformers

Transformers have become the default architecture for DRL: - **Attention mechanisms** handle long-range dependencies in partial observability - **Transformer world models**

(TransDreameer) outperform RNN-based counterparts - **Action chunking transformers** (ACT) for robot imitation learning

5.7 Sim-to-Real Transfer

Domain Randomization (varying physics, textures, lighting) remains the dominant method. Recent advances: - **Adaptive DR**: Automatically adjust randomization based on sim-real performance gap - **System Identification**: Estimate real-world parameters online and adapt simulated training accordingly - **Privileged Information**: Train with privileged teacher (knows hidden state) and distill to student (uses only observations)

Part 6 — Comparative Analysis

Algorithm Comparison Table

Method	Core Idea	Strength	Weakness	Sample Efficiency	Real-world Usage
DQN	Q-learning + deep NN + replay buffer	Simple, well-understood	Discrete only; overestimates Q; unstable	Low	Atari, discrete games
PPO	Clipped surrogate policy gradient	Stable; easy to tune; RLHF standard	On-policy; poor sample efficiency	Medium	ChatGPT/Claude alignment, robotics
SAC	Maximum entropy off-policy actor-critic	SOTA continuous control; auto entropy tuning	Complex; slower than PPO	High	Robot manipulation, LunarLander
MuZero	Learned world model + MCTS	No prior game knowledge; strong planning	Very high compute; complex	Very High (planning)	Board games, Atari
Decision Transformer	Offline RL as sequence modeling	Leverages transformer scale; strong offline	Cannot improve beyond data; no online adaptation	N/A (offline)	Offline control, demonstrations

Key Takeaways

1. **Discrete vs. Continuous Action Space:** DQN is the go-to for discrete spaces; SAC/TD3 for continuous. Our experiments confirm DQN cannot handle LunarLanderContinuous (action space $[-1,1]^2$).
2. **On-policy vs. Off-policy:** PPO (on-policy) is more stable but requires more environment interactions. SAC (off-policy) is more sample efficient but more complex to implement correctly.
3. **Sample Efficiency vs. Stability:** There is a fundamental trade-off. Off-policy methods (SAC, DQN) achieve higher sample efficiency but with more potential for instability. On-policy methods (PPO) are more stable but waste data.
4. **Planning vs. Model-Free:** MuZero's planning advantage is most visible in games with long-term consequences (Go, Chess). In reactive tasks (Atari breakout, continuous control), model-free methods are competitive with much less compute.

Part 7 — GitHub / Open Source Ecosystem Survey

7.1 Stable-Baselines3

Repository: `DLR-RM/stable-baselines3`

Stars: 10,000+

Main Purpose: Clean, reliable PyTorch implementations of RL algorithms for research

Supported Algorithms: DQN, PPO, SAC, TD3, A2C, DDPG, HER

Strengths: Well-documented; reliable hyperparameter defaults; active maintenance; broad community

Weaknesses: No distributed training; limited customization for complex architectures

Applications: Research benchmarking, education, quick prototyping

Used extensively in this homework (all 5 bonus experiments).

7.2 RLlib (Ray)

Repository: `ray-project/ray` (submodule: `rllib`)

Stars: 35,000+ (Ray repo)

Main Purpose: Scalable, production-grade distributed RL

Algorithms: 30+ including PPO, SAC, DQN, APPO, IMPALA, MADDPG

Strengths: Seamless multi-GPU/multi-node scaling; multi-agent support; production-tested

Weaknesses: Complex configuration; heavy dependencies

Applications: Large-scale research, industrial RL (recommendation systems, supply chain)

7.3 CleanRL

Repository: `vwxyzjn/cleanrl`

Stars: 5,000+

Main Purpose: Single-file RL implementations for maximum clarity and reproducibility

Philosophy: Each algorithm in a single Python file with minimal abstraction — easy to understand and modify

Strengths: Extremely readable; Weights & Biases integration; reproducibility focus

Weaknesses: Not designed for production use; limited modularity

Applications: Education, research understanding, algorithm prototyping

7.4 OpenAI Spinning Up

Repository: `openai/spinningup`

Stars: 10,000+

Main Purpose: Educational DRL with clean implementations and detailed documentation

Algorithms: VPG, TRPO, PPO, DDPG, TD3, SAC

Strengths: Best-in-class documentation; pedagogical design; math-code alignment

Weaknesses: Limited maintenance (2020 last major update); TensorFlow 1.x for some implementations

Applications: Learning DRL; implementing baseline algorithms for research

7.5 Isaac Gym / Isaac Lab

Repository: `isaac-sim/IsaacLab`

Main Purpose: GPU-accelerated robot learning simulation

Strengths: Unprecedented simulation throughput (thousands of parallel envs/GPU); direct PyTorch integration

Weaknesses: NVIDIA GPU required; complex setup; steep learning curve

Applications: Humanoid locomotion, dexterous manipulation, sim-to-real transfer

7.6 CARLA

Repository: `carla-simulator/carla`

Stars: 12,000+

Main Purpose: Open-source autonomous driving simulator

Strengths: Photorealistic; comprehensive sensor suite; active CARLA Challenge competition

Weaknesses: High compute requirements; slow per-step (Unreal Engine overhead)

Applications: Self-driving research, DRL-based planning, computer vision for driving

7.7 Habitat Lab

Repository: `facebookresearch/habitat-lab`

Stars: 2,000+

Main Purpose: Embodied AI research in photorealistic 3D environments

Strengths: Largest indoor 3D scene datasets; strong sim-to-real transfer results

Weaknesses: Primarily navigation tasks; less suitable for manipulation

Applications: Indoor navigation, social robotics, visual question answering in 3D

7.8 FinRL

Repository: `AI4Finance-Foundation/FinRL`

Stars: 10,000+

Main Purpose: DRL for quantitative finance

Strengths: End-to-end financial RL pipeline; multiple data sources; paper reproductions

Weaknesses: Research-grade (not production-ready trading); market impact not modeled

Applications: Algorithmic trading research, portfolio optimization benchmarking

7.9 MineDojo

Repository: `MineDojo/MineDojo`

Stars: 1,500+

Main Purpose: Open-ended embodied agent learning in Minecraft

Strengths: Thousands of diverse tasks; natural language task specification; internet-scale training data (YouTube)

Weaknesses: Complex setup; high compute for pretraining

Applications: Embodied AI research, open-ended learning, vision-language-action models

7.10 Unity ML-Agents

Repository: `Unity-Technologies/ml-agents`

Stars: 17,000+

Main Purpose: RL and imitation learning in Unity game environments

Algorithms: PPO, SAC (built-in); any external algorithm via Python API

Strengths: Visual environment design; cross-platform; large community

Applications: Game AI, robotics simulation, behavior design

Experimental Results Summary

Bonus Experiment Overview

Experiment	Algorithm	Environment	Key Result
Bonus 1	DQN vs PPO	CartPole-v1	PPO: 500.0; DQN: 144.4 (catastrophic forgetting at ep.~300)
Bonus 2	PPO (3 envs)	CartPole, MountainCar, Acrobot	CartPole: 475; Acrobot: -100; MountainCar: -120
Bonus 3	PPO ablation	CartPole-v1	Best: lr=3e-4, clip=0.1; lr=1e-4 unstable
Bonus 4	DQN	10×10 GridWorld	22-step optimal path; 99% success rate; Reward: 98
Bonus 5	SAC	LunarLanderContinuous-v3	271.7 ± 15.5  (300k steps, above 200 threshold)

Key Findings

- PPO vs DQN on CartPole:** PPO's on-policy stability advantage is clear — it converges monotonically to 500 while DQN suffers from catastrophic forgetting after discovering good policies early. This mirrors RLHF practice where PPO's stability is crucial.
- Sparse rewards are hard:** MountainCar (reward only at goal) challenges on-policy PPO without explicit exploration mechanisms. This motivates entropy-based exploration (SAC) or curiosity-driven intrinsic rewards (ICM).
- SAC enables continuous control:** LunarLanderContinuous-v3 is fundamentally inaccessible to DQN (continuous $[-1,1]^2$ action space). SAC's maximum entropy framework achieves 271.7 reward (well above the 200 threshold), with entropy auto-tuning (α : $\sim 1.0 \rightarrow \sim 0.2$) visibly verified — demonstrating the necessity of continuous control methods for real-world robotics.
- Hyperparameter sensitivity:** Bonus 3 ablation shows that lr=3e-4 (PPO default) with clip=0.1 is remarkably robust, while lr=1e-4 with large clip_range causes instability. This validates the conventional PPO hyperparameter wisdom.

Conclusion

Deep Reinforcement Learning has matured from a research curiosity (Atari-playing DQN, 2015) to a foundational technology for real-world AI systems. Key observations:

1. **Algorithm selection is task-dependent:** DQN for discrete tasks; SAC/TD3 for continuous; PPO for stable, scalable training including RLHF.
2. **Agentic AI is the frontier:** Combining DRL (for action generation and reward optimization) with LLMs (for planning and reasoning) and world models (for imagination-based planning) represents the most promising path toward general autonomous agents.
3. **Data and compute are the bottleneck:** World-model pretraining, offline RL, and foundation models are strategies to reduce the data requirements that limit real-world DRL deployment.
4. **Safety and alignment are unsolved:** Constrained RL, RLHF/RLAIF, and Constitutional AI are active research areas. No deployed system fully solves the alignment problem — an important ethical consideration for all DRL practitioners.

References

1. Mnih, V., et al. (2015). Human-level control through deep reinforcement learning. *Nature*, 518, 529–533.
2. van Hasselt, H., Guez, A., & Silver, D. (2016). Deep reinforcement learning with double Q-learning. *AAAI*.
3. Schulman, J., et al. (2017). Proximal Policy Optimization Algorithms. *arXiv:1707.06347*.
4. Haarnoja, T., et al. (2018). Soft Actor-Critic: Off-Policy Maximum Entropy Deep Reinforcement Learning. *ICML*.
5. Fujimoto, S., et al. (2018). Addressing Function Approximation Error in Actor-Critic Methods. *ICML*.
6. Schrittwieser, J., et al. (2020). Mastering Atari, Go, Chess and Shogi by Planning with a Learned Model. *Nature*, 588, 604–609.
7. Chen, L., et al. (2021). Decision Transformer: Reinforcement Learning via Sequence Modeling. *NeurIPS*.
8. Vinyals, O., et al. (2019). Grandmaster level in StarCraft II using multi-agent reinforcement learning. *Nature*, 575, 350–354.
9. Bai, Y., et al. (2022). Constitutional AI: Harmlessness from AI Feedback. *arXiv:2212.06074*.
10. Hafner, D., et al. (2023). Mastering Diverse Domains through World Models. *arXiv:2301.04104* (DreamerV3).
11. Chi, C., et al. (2023). Diffusion Policy: Visuomotor Policy Learning via Action Diffusion. *RSS*.
12. Reed, S., et al. (2022). A Generalist Agent (Gato). *arXiv:2205.06175*.
13. Wang, G., et al. (2023). Voyager: An Open-Ended Embodied Agent with Large Language Models. *NeurIPS*.

14. Liu, X., et al. (2022). FinRL: A Deep Reinforcement Learning Library for Automated Stock Trading. *ACM ICAIF*.
15. Raffin, A., et al. (2021). Stable-Baselines3: Reliable Reinforcement Learning Implementations. *JMLR*.
16. Chitta, K., et al. (2022). TransFuser: Imitation with Transformer-Based Sensor Fusion. *TPAMI*.
17. Kim, M., et al. (2024). OpenVLA: An Open-Source Vision-Language-Action Model. *arXiv:2406.09246*.